

Aim:

To read stock market data into a DataFrame, use window functions to calculate the moving average price for each stock, and display the results.

Theory:

Apache Spark is a fast and general-purpose distributed computing system that provides APIs for large-scale data processing. Spark DataFrames allow users to manipulate structured data using a domain-specific language similar to SQL.

Window functions in Spark enable operations like running totals, ranking, and moving averages across a set of related rows.

Steps:

Step 1 : Download & extract Spark

```
wget https://archive.apache.org/dist/spark/spark-3.5.0/spark-3.5.0-bin-hadoop3.tgz
```

```
devashree@DEVASHREE:~$ wget https://archive.apache.org/dist/spark/spark-3.5.0/spark-3.5.0-bin-hadoop3.tgz
--2025-10-23 06:37:31-- https://archive.apache.org/dist/spark/spark-3.5.0/spark-3.5.0-bin-hadoop3.tgz
Resolving archive.apache.org (archive.apache.org)... 65.108.204.189, 2a01:4f9:1:a:a084::2
Connecting to archive.apache.org (archive.apache.org)|65.108.204.189|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 400395283 (382M) [application/x-gzip]
Saving to: 'spark-3.5.0-bin-hadoop3.tgz'

spark-3.5.0-bin-hadoop3.tgz      100%[=====] 381.85M  2.45MB/s   in 5m 44s

2025-10-23 06:43:36 (1.11 MB/s) - 'spark-3.5.0-bin-hadoop3.tgz' saved [400395283/400395283]
```

```
tar -xvzf spark-3.5.0-bin-hadoop3.tgz
```

```
devashree@DEVASHREE:~$ tar -xvzf spark-3.5.0-bin-hadoop3.tgz
spark-3.5.0-bin-hadoop3/
spark-3.5.0-bin-hadoop3/kubernetes/
spark-3.5.0-bin-hadoop3/kubernetes/tests/
spark-3.5.0-bin-hadoop3/kubernetes/tests/pyfiles.py
spark-3.5.0-bin-hadoop3/kubernetes/tests/decommissioning.py
spark-3.5.0-bin-hadoop3/kubernetes/tests/autoscale.py
spark-3.5.0-bin-hadoop3/kubernetes/tests/python_executable_check.py
spark-3.5.0-bin-hadoop3/kubernetes/tests/worker_memory_check.py
spark-3.5.0-bin-hadoop3/kubernetes/tests/py_container_checks.py
spark-3.5.0-bin-hadoop3/kubernetes/tests/decommissioning_cleanup.py
spark-3.5.0-bin-hadoop3/kubernetes/dockerfiles/
spark-3.5.0-bin-hadoop3/kubernetes/dockerfiles/spark/
```

Step 2 : Create sample CSV (stock_prices.csv)

Date	Stock	Close
2025-08-01	AAPL	180
2025-08-02	AAPL	182
2025-08-03	AAPL	181
2025-08-04	AAPL	185
2025-08-01	TSLA	210
2025-08-02	TSLA	215
2025-08-03	TSLA	220
2025-08-04	TSLA	205
2025-08-01	MSFT	320

2025-08-02,MSFT,325
2025-08-03,MSFT,322
2025-08-04,MSFT,330

```
devashree@DEVASHREE:~$ cat > ~/stock_prices.csv <<'EOL'
Date,Stock,Close
2025-08-01,AAPL,180
2025-08-02,AAPL,182
2025-08-03,AAPL,181
2025-08-04,AAPL,185
2025-08-01,TSLA,210
2025-08-02,TSLA,215
2025-08-03,TSLA,220
2025-08-04,TSLA,205
2025-08-01,MSFT,320
2025-08-02,MSFT,325
2025-08-03,MSFT,322
2025-08-04,MSFT,330
EOL
```

Step 3 : Upload CSV to HDFS

```
hdfs dfs -mkdir -p /user/hduser/stocks
```

```
devashree@DEVASHREE:~$ hadoop dfs -mkdir -p /user/hduser/stocks  
2025-10-23 06:50:10,870 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

```
hdfs dfs -put ~/stock_prices.csv /user/hduser/stocks/
```

```
devashree@DEVASHREE:~$ hdfs dfs -put ~/stock_prices.csv /user/hduser/stocks/
2025-10-23 06:50:40.015 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

```
hdfs dfs -ls /user/hduser/stocks
```

```
devashree@DEVAshree:~$ hdfs dfs -ls /user/hduser/stocks
2025-10-23 06:50:50,248 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
 1 -rw-r--r-- 1 hduser supergroup          128 2025-10-23 06:50 /user/hduser/stocks/stock_prices.csv
```

Step 4 : Launch spark-shell (Scala)

```
spark-shell --master local[*]
```

Step 5 : Load the data in Spark

```
val stocksDF = spark.read  
.option("header", "true")  
.option("inferSchema", "true")  
.csv("file:///home/devashree/stock_prices.csv")  
  
scala> val stocksDF = spark.read .option("header", "true") .option("inferSchema", "true") .csv("file:///home/devashree/  
/stock_prices.csv")  
stocksDF: org.apache.spark.sql.DataFrame = [Date: date, Stock: string ... 1 more field]
```

stocksDF.show()

```
scala> stocksDF.show()  
+-----+-----+-----+  
| Date | Stock | Close |  
+-----+-----+-----+  
| 2025-08-01 | AAPL | 180 |  
| 2025-08-02 | AAPL | 182 |  
| 2025-08-03 | AAPL | 181 |  
| 2025-08-04 | AAPL | 185 |  
| 2025-08-01 | TSLA | 210 |  
| 2025-08-02 | TSLA | 215 |  
| 2025-08-03 | TSLA | 220 |  
| 2025-08-04 | TSLA | 205 |  
| 2025-08-01 | MSFT | 320 |  
| 2025-08-02 | MSFT | 325 |  
| 2025-08-03 | MSFT | 322 |  
| 2025-08-04 | MSFT | 330 |  
+-----+-----+-----+
```

stocksDF.printSchema()

```
scala> stocksDF.printSchema()  
root  
|-- Date: date (nullable = true)  
|-- Stock: string (nullable = true)  
|-- Close: integer (nullable = true)
```

Step 6 : Compute 3-day moving average

```
import org.apache.spark.sql.expressions.Window  
import org.apache.spark.sql.functions._  
  
val windowSpec = Window.partitionBy("Stock").orderBy("Date").rowsBetween(-2, 0)  
  
val resultDF = stocksDF  
.withColumn("MA_3", round(avg(col("Close")).over(windowSpec), 2))  
.select("Date", "Stock", "Close", "MA_3")  
.orderBy("Stock", "Date")  
  
resultDF.show(50, false)
```

```

scala> import org.apache.spark.sql.expressions.Window
import org.apache.spark.sql.expressions.Window

scala> import org.apache.spark.sql.functions._
import org.apache.spark.sql.functions._

scala> val windowSpec = Window.partitionBy("Stock").orderBy("Date").rowsBetween(-2, 0)
windowSpec: org.apache.spark.sql.expressions.WindowSpec = org.apache.spark.sql.expressions.WindowSpec@76e13660

scala> val movingAvgDF = stocksDF.withColumn("MA_3", round(avg("Close").over(windowSpec), 2))
movingAvgDF: org.apache.spark.sql.DataFrame = [Date: date, Stock: string ... 2 more fields]

scala> val resultDF = movingAvgDF.select("Date", "Stock", "Close", "MA_3").orderBy("Stock", "Date")
resultDF: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [Date: date, Stock: string ... 2 more fields]

scala> resultDF.show(50, false)
+-----+---+-----+
|Date   |Stock|Close|MA_3 |
+-----+---+-----+
|2025-08-01|AAPL |180  |180.0 |
|2025-08-02|AAPL |182  |181.0 |
|2025-08-03|AAPL |181  |181.0 |
|2025-08-04|AAPL |185  |182.67|
|2025-08-01|MSFT |320  |320.0 |
|2025-08-02|MSFT |325  |322.5 |
|2025-08-03|MSFT |322  |322.33|
|2025-08-04|MSFT |330  |325.67|
|2025-08-01|TSLA |210  |210.0 |
|2025-08-02|TSLA |215  |212.5 |
|2025-08-03|TSLA |220  |215.0 |
|2025-08-04|TSLA |205  |213.33|
+-----+---+-----+

```

EXERCISE:

- 1) Load the stock data and display the first 10 rows of the DataFrame.

```
stocksDF.show(10, false)
```

```

scala> stocksDF.show(10, false)
+-----+---+-----+
|Date   |Stock|Close|
+-----+---+-----+
|2025-08-01|AAPL |180  |
|2025-08-02|AAPL |182  |
|2025-08-03|AAPL |181  |
|2025-08-04|AAPL |185  |
|2025-08-01|TSLA |210  |
|2025-08-02|TSLA |215  |
|2025-08-03|TSLA |220  |
|2025-08-04|TSLA |205  |
|2025-08-01|MSFT |320  |
|2025-08-02|MSFT |325  |
+-----+---+-----+
only showing top 10 rows

```

- 2) Filter the DataFrame to show only the rows where the stock symbol is AAPL.

```
val aaplDF = stocksDF.filter(col("Stock") === "AAPL")
aaplDF.show(false)
```

```

scala> val aaplDF = stocksDF.filter(col("Stock") === "AAPL")
aaplDF: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [Date: date, Stock: string ... 1 more field]

scala> aaplDF.show(false)
+-----+---+---+
|Date    |Stock|Close|
+-----+---+---+
|2025-08-01|AAPL |180  |
|2025-08-02|AAPL |182  |
|2025-08-03|AAPL |181  |
|2025-08-04|AAPL |185  |
+-----+---+---+

```

- 3) Sort the data by the Close price in descending order and display the top 5 rows.

```
stocksDF.orderBy(col("Close").desc).show(5, false)
```

```

scala> stocksDF.orderBy(col("Close").desc).show(5, false)
+-----+---+---+
|Date    |Stock|Close|
+-----+---+---+
|2025-08-04|MSFT |330  |
|2025-08-02|MSFT |325  |
|2025-08-03|MSFT |322  |
|2025-08-01|MSFT |320  |
|2025-08-03|TSLA |220  |
+-----+---+---+
only showing top 5 rows

```

- 4) Count the total number of records in the stock DataFrame.

```
println("Total records: " + stocksDF.count())
```

```

scala> println("Total number of records: " + stocksDF.count())
Total number of records: 12

```

- 5) Save the filtered data for TSLA into a separate directory in HDFS.

```
val tslaRead=spark.read.option("header","true").csv("file:///home/devashree/tsla_filtered")
```

```
tslaRead.show()
```

```

scala> val tslaRead = spark.read.option("header", "true").csv("file:///home/devashree/tsla_filtered")
tslaRead: org.apache.spark.sql.DataFrame = [Date: string, Stock: string ... 1 more field]

scala> tslaRead.show()
+-----+---+---+
|      Date|Stock|Close|
+-----+---+---+
|2025-08-01| TSLA| 210|
|2025-08-02| TSLA| 215|
|2025-08-03| TSLA| 220|
|2025-08-04| TSLA| 205|
+-----+---+---+

```

6) Display only the Date and Close columns for all stocks.

```
stocksDF.select("Date", "Close").show(20, false)
```

```
scala> stocksDF.select("Date", "Close").show(20, false)
+-----+----+
|Date      |Close|
+-----+----+
|2025-08-01|180  |
|2025-08-02|182  |
|2025-08-03|181  |
|2025-08-04|185  |
|2025-08-01|210  |
|2025-08-02|215  |
|2025-08-03|220  |
|2025-08-04|205  |
|2025-08-01|320  |
|2025-08-02|325  |
|2025-08-03|322  |
|2025-08-04|330  |
+-----+----+
```

7) Find the unique stock symbols present in the dataset.

```
stocksDF.select("Stock").distinct().show(false)
```

```
scala> stocksDF.select("Stock").distinct().show(false)
+---+
|Stock|
+---+
|AAPL  |
|TSLA  |
|MSFT  |
+---+
```

8) Display only the records where the Close price is greater than 200.

```
stocksDF.filter(col("Close") > 200).show(false)
```

```
scala> stocksDF.filter(col("Close") > 200).show(false)
+-----+----+----+
|Date      |Stock|Close|
+-----+----+----+
|2025-08-01|TSLA |210  |
|2025-08-02|TSLA |215  |
|2025-08-03|TSLA |220  |
|2025-08-04|TSLA |205  |
|2025-08-01|MSFT |320  |
|2025-08-02|MSFT |325  |
|2025-08-03|MSFT |322  |
|2025-08-04|MSFT |330  |
+-----+----+----+
```

Preparation	20	
Implementation	20	
Viva	15	
Output	10	
Record	10	
Total	75	

Result:

Thus, the reading of stock market data into a DataFrame, use window functions to calculate the moving average price for each stock, and display the results.

Aim:

To Set up a version control repository (e.g., GitHub), configure a CI tool (e.g., Jenkins, GitLab CI), define stages for building, testing, and deploying the application, and trigger the pipeline on code commits.

Steps to be Implemented:**1) Create a Maven project (Hello World)****Open a terminal and run:**

```
cd ~
```

```
mvn archetype:generate -DgroupId=com.example \
```

```
-DartifactId=firstProject \
-DarchetypeArtifactId=maven-archetype-quickstart \
-DinteractiveMode=false
```

```
cd firstProject
```

```
>mvn archetype:generate -DgroupId=com.example -DartifactId=firstProject -DarchetypeArtifactId=maven-archetype-quickstart
-DinteractiveMode=false
[INFO] Scanning for projects...
Downloading from central: https://repo.maven.apache.org/maven2/org/apache/maven/plugins/maven-archetype-plugin/maven-metadata.xml
Downloaded from central: https://repo.maven.apache.org/maven2/org/apache/maven/plugins/maven-archetype-plugin/maven-metadata.xml (1.1 kB at 1.7 kB/s)
[INFO]
[INFO] -----< org.apache.maven:standalone-pom >-----
[INFO] Building Maven Stub Project (No POM) 1
[INFO] -----[ pom ]-----
[INFO]
[INFO] >>> archetype:3.4.1:generate (default-cli) > generate-sources @ standalone-pom >>>
[INFO]
[INFO] <<< archetype:3.4.1:generate (default-cli) < generate-sources @ standalone-pom <<<
[INFO]
[INFO]
[INFO] --- archetype:3.4.1:generate (default-cli) @ standalone-pom ---
[INFO] Generating project in Batch mode
Downloading from central: https://repo.maven.apache.org/maven2/archetype-catalog.xml
Downloaded from central: https://repo.maven.apache.org/maven2/archetype-catalog.xml (17 MB at 28 MB/s)
```

Verify: You should see a pom.xml and a src folder.

```
$ cd firstProject
$ ls
pom.xml  src
```

Set JDK version in pom.xml:

```
<properties>
    <maven.compiler.source>11</maven.compiler.source>
    <maven.compiler.target>11</maven.compiler.target>
</properties>
```

```

GNU nano 7.2                               pom.xml *
<project xmlns="http://maven.apache.org/POM/4.0.0" xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://maven.apache.org/POM/4.0.0 http://maven.apache.org/maven-v4_0_0.xsd">
  <modelVersion>4.0.0</modelVersion>
  <groupId>com.example</groupId>
  <artifactId>firstProject</artifactId>
  <packaging>jar</packaging>
  <version>1.0-SNAPSHOT</version>
  <name>firstProject</name>
  <url>http://maven.apache.org</url>
  <dependencies>
    <dependency>
      <groupId>junit</groupId>
      <artifactId>junit</artifactId>
      <version>3.8.1</version>
      <scope>test</scope>
    </dependency>
  </dependencies>
  <properties>
    <maven.compiler.source>11</maven.compiler.source>
    <maven.compiler.target>11</maven.compiler.target>
  </properties>
</project>

```

Verify: File compiles later via Maven without errors.

Build locally: mvn clean install

Verify: Look for BUILD SUCCESS and a JAR at target/firstProject-1.0-SNAPSHOT.jar.

```

$ mvn clean install
[INFO] Scanning for projects...
[INFO]
[INFO] -----< com.example:firstProject >-----
[INFO] Building firstProject 1.0-SNAPSHOT
[INFO] -----[ jar ]-----
[INFO]
[INFO] --- maven-clean-plugin:2.5:clean (default-clean) @ firstProject ---
[INFO]
[INFO] --- maven-resources-plugin:2.6:resources (default-resources) @ firstProject ---
[WARNING] Using platform encoding (UTF-8 actually) to copy filtered resources, i.e. build is platform dependent!
[INFO] skip non existing resourceDirectory /home/devashree/firstProject/src/main/resources
[INFO]
[INFO] --- maven-compiler-plugin:3.1:compile (default-compile) @ firstProject ---
[INFO] Changes detected - recompiling the module!
[WARNING] File encoding has not been set, using platform encoding UTF-8, i.e. build is platform dependent!
[INFO] Compiling 1 source file to /home/devashree/firstProject/target/classes
[INFO]
[INFO] --- maven-resources-plugin:2.6:testResources (default-testResources) @ firstProject ---
[WARNING] Using platform encoding (UTF-8 actually) to copy filtered resources, i.e. build is platform dependent!
[INFO] skip non existing resourceDirectory /home/devashree/firstProject/src/test/resources
[INFO]
[INFO] --- maven-compiler-plugin:3.1:testCompile (default-testCompile) @ firstProject ---
[INFO] Changes detected - recompiling the module!
[WARNING] File encoding has not been set, using platform encoding UTF-8, i.e. build is platform dependent!
[INFO] Compiling 1 source file to /home/devashree/firstProject/target/test-classes

```

2) Initialize Git and make your first commit

Configure your Git identity (one-time):

```

git config --global user.name "Your Name"
git config --global user.email "you@example.com"

```

Initialize and commit:

```

git init
git add .
git commit -m "Initial commit"

```

```

$ git config --global user.name "Shree2722"
$ git config --global user.email "shree45343@gmail.com"
$ git init
hint: Using 'master' as the name for the initial branch. This default branch name
hint: is subject to change. To configure the initial branch name to use in all
hint: of your new repositories, which will suppress this warning, call:
hint:
hint:   git config --global init.defaultBranch <name>
hint:
hint: Names commonly chosen instead of 'master' are 'main', 'trunk' and
hint: 'development'. The just-created branch can be renamed via this command:
hint:
hint:   git branch -m <name>
Initialized empty Git repository in /home/devashree/firstProject/.git/
$ git add .
$ git commit -m "Initial commit"
[master (root-commit) f4960b4] Initial commit
 13 files changed, 151 insertions(+)
create mode 100644 pom.xml
create mode 100644 src/main/java/com/example/App.java
create mode 100644 src/test/java/com/example/AppTest.java
create mode 100644 target/classes/com/example/App.class
create mode 100644 target/firstProject-1.0-SNAPSHOT.jar
create mode 100644 target/maven-archiver/pom.properties
create mode 100644 target/maven-status/maven-compiler-plugin/compile/default-compile/createdFiles.lst
create mode 100644 target/maven-status/maven-compiler-plugin/compile/default-compile/inputFiles.lst
create mode 100644 target/maven-status/maven-compiler-plugin/testCompile/default-testCompile/createdFiles.lst
create mode 100644 target/maven-status/maven-compiler-plugin/testCompile/default-testCompile/inputFiles.lst
create mode 100644 target/surefire-reports/TEST-com.example.AppTest.xml
create mode 100644 target/surefire-reports/com.example.AppTest.txt
create mode 100644 target/test-classes/com/example/AppTest.class

```

3) Push the project to GitHub

Create an empty GitHub repository named firstProject (no README).

Create a new repository

Repositories contain a project's files and version history. Have a project elsewhere? [Import a repository](#).

Required fields are marked with an asterisk (*).

1 General

Owner * Shree2722 / **Repository name *** firstProject firstProject is available.

Great repository names are short and memorable. How about `symmetrical-couscous?`

Description

0 / 350 characters

2 Configuration

Choose visibility * Public

Choose who can see and commit to this repository

Add README READMEs can be used as longer descriptions. [About READMEs](#) Off

Add .gitignore .gitignore tells git which files not to track. [About ignoring files](#) No .gitignore

Add license Licenses explain how others can use your code. [About licenses](#) No license

Create repository

Connect and push:

```

git remote add origin https://github.com/<your-username>/firstProject.git
git branch -M main
git push -u origin main

```

```

$ git remote add origin https://github.com/Shree2722/firstProject.git
$ git branch -M main

$ git push -u origin main
Username for 'https://github.com': Shree2722
Password for 'https://Shree2722@github.com':
Enumerating objects: 39, done.
Counting objects: 100% (39/39), done.
Delta compression using up to 12 threads
Compressing objects: 100% (19/19), done.
Writing objects: 100% (39/39), 6.03 KiB | 882.00 KiB/s, done.
Total 39 (delta 0), reused 0 (delta 0), pack-reused 0
To https://github.com/Shree2722/firstProject.git
 * [new branch]      main -> main
branch 'main' set up to track 'origin/main'.
$ |

```

4) Install Jenkins (on Linux)

sudo apt update

sudo apt install openjdk-11-jdk -y

wget -q -O - https://pkg.jenkins.io/debian-stable/jenkins.io.key | sudo apt-key add -

sudo sh -c 'echo deb https://pkg.jenkins.io/debian-stable binary/ > /etc/apt/sources.list.d/jenkins.list'

sudo apt update

sudo apt install jenkins -y

sudo systemctl start jenkins

sudo systemctl enable jenkins

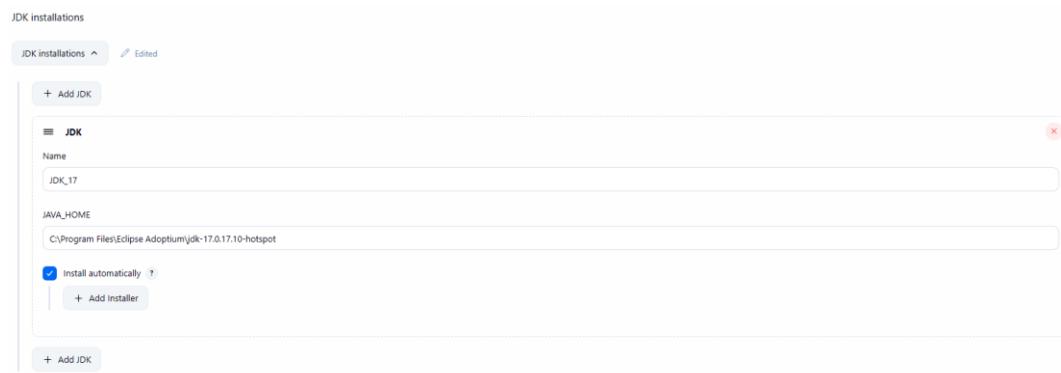
Verify:

- Visit <http://<server-ip>:8080>
- Unlock Jenkins with the initial admin password (/var/lib/jenkins/secrets/initialAdminPassword).

5) Configure tools and plugins in Jenkins

1. In Jenkins: Manage Jenkins → Global Tool Configuration

- Add JDK 11 (or make sure “Install automatically” is set).
- Add Maven (install automatically or specify a path).



2. Install plugins: **Git** and **Maven Integration** (via **Manage Jenkins → Plugins**).
Verify: You can see JDK and Maven listed under Global Tool Configuration without warnings.

6) Create a Jenkins Maven job

1. New Item → Maven project

Name: firstProject

2. Source Code Management → Git

- Repository URL: <https://github.com/<your-username>/firstProject.git>
- Credentials: Add your GitHub credentials if the repo is private
- Branch: main

Git ?
Repositories ?

Repository URL ?
`https://github.com/Shree2722/firstProject.git`

Credentials ?
- none -

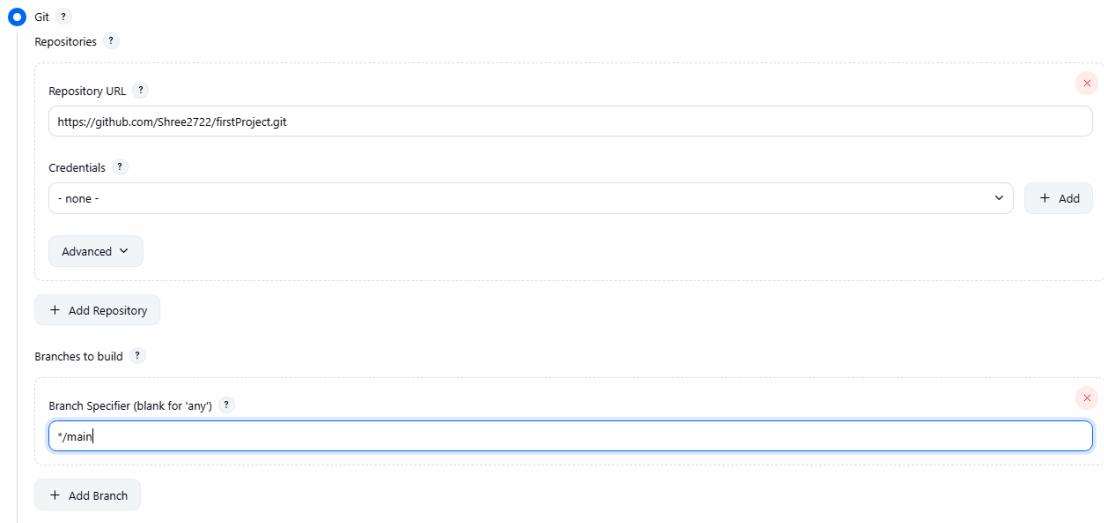
Advanced ▾

+ Add Repository

Branches to build ?

Branch Specifier (blank for 'any') ?
`*/main`

+ Add Branch



3. Build

- o Root POM: pom.xml
- o Goals: clean install

4. Save.

Verify: Job shows your repo URL and branch correctly.

≡ Invoke top-level Maven targets ?

Maven Version
(Default)

Goals
clean install

Advanced ^ 

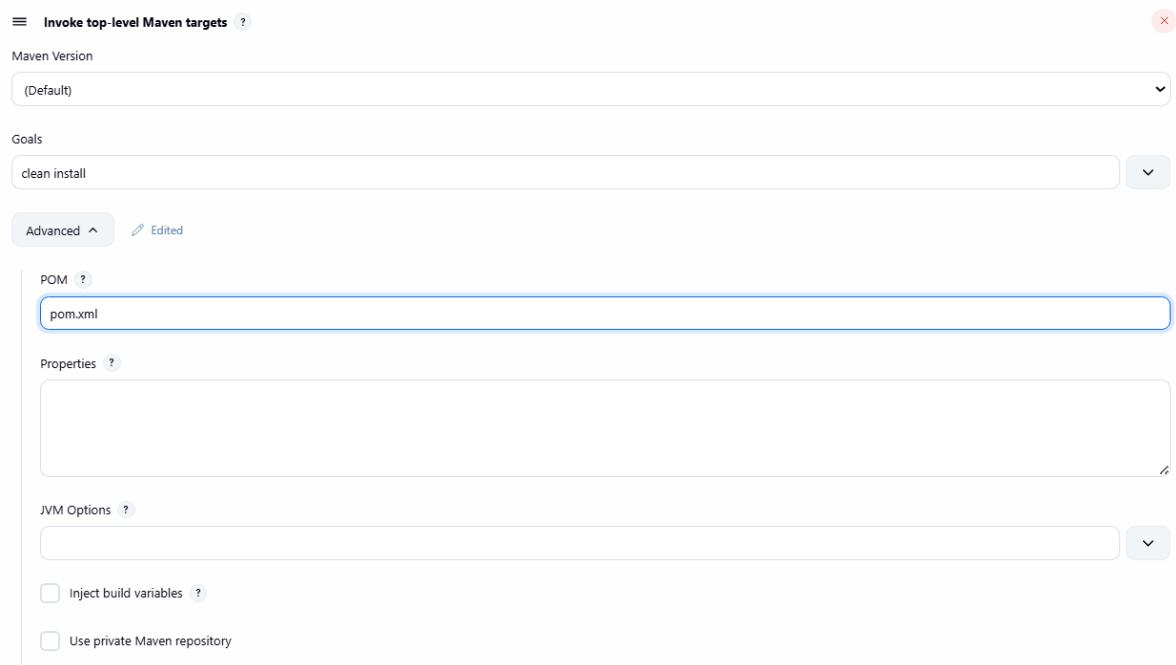
POM ?
pom.xml

Properties ?

JVM Options ?

Inject build variables ?

Use private Maven repository



7) Run the pipeline and confirm artifacts

1. Click **Build Now**.

2. Open **Console Output**.

Verify: You should see BUILD SUCCESS. Also check the workspace:

3. <JENKINS_HOME>/workspace/firstProject/target/firstProject-1.0-SNAPSHOT.jar

Jenkins | firstProject | #1 | Console Output

```
[INFO] --- compiler:3.13.0:testCompile (default-testCompile) @ firstProject ---
[INFO] Recompiling the module because of changed dependency.
[WARNING] File encoding has not been set, using platform encoding UTF-8, i.e. build is platform dependent!
[INFO] Compiling 1 source file with javac [debug target 11] to target/test-classes
[WARNING] system.modules path not set in conjunction with -source 11
[INFO]
[INFO] --- surefire:3.2.5:test (default-test) @ firstProject ---
[INFO] Using auto detected provider org.apache.maven.surefire:junit4Provider
[INFO]
[INFO] -----
[INFO] T E S T S
[INFO] -----
[INFO] Running com.example.AppTest
[INFO] Tests run: 1, Failures: 0, Errors: 0, Skipped: 0, Time elapsed: 0.084 s -- in com.example.AppTest
[INFO]
[INFO] Results:
[INFO]
[INFO] Tests run: 1, Failures: 0, Errors: 0, Skipped: 0
[INFO]
[INFO] --- jar:3.4.1:jar (default-jar) @ firstProject ---
[INFO] Building jar: D:\Jenkins_home\workspace\firstProject\target\firstProject-1.0-SNAPSHOT.jar
[INFO]
[INFO] --- install:3.1.2:install (default-install) @ firstProject ---
[INFO] Installing D:\Jenkins_home\workspace\firstProject\pom.xml to C:\Users\VH.DevaShree\.m2\repository\com\example\firstProject\1.0-SNAPSHOT\firstProject-1.0-SNAPSHOT.pom
[INFO] Installing D:\Jenkins_home\workspace\firstProject\target\firstProject-1.0-SNAPSHOT.jar to C:\Users\VH.DevaShree\.m2\repository\com\example\firstProject\1.0-SNAPSHOT\firstProject-1.0-SNAPSHOT.jar
[INFO]
[INFO] -----
[INFO] BUILD SUCCESS
[INFO] -----
[INFO] Total time: 11.497 s
[INFO] Finished at: 2025-11-05T16:21:38+05:30
[INFO] -----
[INFO] Finished: SUCCESS
```

This PC > New Volume (D:) > Jenkins_home > workspace > firstProject > target >

Name	Date modified	Type	Size
classes	11/5/2025 4:21 PM	File folder	
generated-sources	11/5/2025 4:21 PM	File folder	
generated-test-sources	11/5/2025 4:21 PM	File folder	
maven-archiver	11/5/2025 4:21 PM	File folder	
maven-status	11/5/2025 4:21 PM	File folder	
surefire-reports	11/5/2025 4:21 PM	File folder	
test-classes	11/5/2025 4:21 PM	File folder	
firstProject-1.0-SNAPSHOT.jar	11/5/2025 4:21 PM	jarfile	3 KB

Preparation	20	
Implementation	20	
Viva	15	
Output	10	
Record	10	
Total	75	

Result:

Thus, the CI pipeline was successfully set up using GitHub and Jenkins, which automatically built and tested the application whenever new code was committed.

Aim:

To write a Dockerfile to package the application into a container, set up a container registry (e.g., Docker Hub, Amazon ECR), and deploy the containerized application to a container orchestration platform (e.g., Kubernetes, Docker Swarm) using the CI/CD pipeline..

Prerequisites:

- **Docker:** A platform to build, package, and run applications in **containers**, which are lightweight, portable, and isolated environments.
- **Kubernetes (K8s):** A container orchestration system that manages deployment, scaling, and networking of containers across clusters.
- **CI/CD:** Continuous Integration and Continuous Deployment pipelines automate building, testing, and deploying code.

Requirements:

1. **Windows 10/11 with WSL2**
 - WSL2 allows Linux-based containers to run on Windows seamlessly.
2. **Docker Desktop**
 - Provides Docker Engine, CLI, and optional Kubernetes cluster for local testing.
3. **Node.js Project**
 - Application to deploy; in this manual, an Express.js app.

Step 1: Install Docker Desktop & Enable Kubernetes

- Docker Desktop includes Docker Engine + Docker CLI + optional Kubernetes cluster.
- Enabling Kubernetes allows you to test container orchestration locally without setting up a full cloud cluster.

Steps:-

1. Install Docker Desktop from docker.com.
2. Settings → Enable WSL2 Integration → Select your distro.
3. Settings → Kubernetes → Enable Kubernetes → Apply & Restart.
4. Verify:

```
docker version  
kubectl version --client  
wsl --list --verbose
```

```
$ docker version
Client:
  Version:          28.5.1
  API version:     1.51
  Go version:      go1.24.8
  Git commit:      e180ab8
  Built:           Wed Oct  8 12:16:30 2025
  OS/Arch:         linux/amd64
  Context:          default

Server: Docker Desktop 4.49.0 (208700)
Engine:
  Version:          28.5.1
  API version:     1.51 (minimum version 1.24)
  Go version:      go1.24.8
  Git commit:      f8215cc
  Built:           Wed Oct  8 12:17:24 2025
  OS/Arch:         linux/amd64
  Experimental:    false
  containerd:
    Version:        1.7.27
    GitCommit:      05044ec0a9a75232cad458027ca83437aae3f4da
  runc:
    Version:        1.2.5
    GitCommit:      v1.2.5-0-g59923ef
  docker-init:
    Version:        0.19.0
    GitCommit:      de40ad0
$ kubectl version --client
Client Version: v1.34.1
Kustomize Version: v5.7.1
```

- docker version - checks Docker Engine and CLI.
- kubectl version - checks Kubernetes client connectivity.
- wsl --list --verbose - confirms WSL2 distro status.

Step 2: Create Node.js Project

- Node.js is a runtime environment for JavaScript outside the browser.
- Express.js is a minimal web framework for Node.js, used for handling HTTP requests.

Run the below commands:-

```
mkdir C:\Users\welcome\Desktop\regapp-node
cd C:\Users\welcome\Desktop\regapp-node
npm init -y
npm install express
```

```

>cd C:\Users\M.DevaShree\Desktop
>mkdir regapp-node
>cd regapp-node
>npm init -y
Wrote to C:\Users\M.DevaShree\Desktop\regapp-node\package.json:

{
  "name": "regapp-node",
  "version": "1.0.0",
  "main": "index.js",
  "scripts": {
    "test": "echo \"Error: no test specified\" && exit 1"
  },
  "keywords": [],
  "author": "",
  "license": "ISC",
  "description": ""
}

>npm install express
added 68 packages, and audited 69 packages in 5s

16 packages are looking for funding
  run `npm fund` for details

found 0 vulnerabilities

```

Create app.js:

```

const express = require('express');
const app = express();
const PORT = process.env.PORT || 8080;

app.get('/', (req, res) => res.send('Hello from Node.js Docker App!'));

app.listen(PORT, () => console.log(`Server running on port ${PORT}`));

```

- `process.env.PORT` → allows Kubernetes or Docker to map container ports dynamically.
- `app.get('/', ...)` → basic HTTP GET route.

Step 3: Dockerize Node.js App

- Dockerfile defines how to build a container image.
- Layers in Docker improve efficiency (cached during builds).

Create Dockerfile:

Paste the above code in the Dockerfile.

```
FROM node:20
WORKDIR /app
COPY package*.json .
RUN npm install
COPY ..
EXPOSE 8080
CMD ["npm", "start"]
```

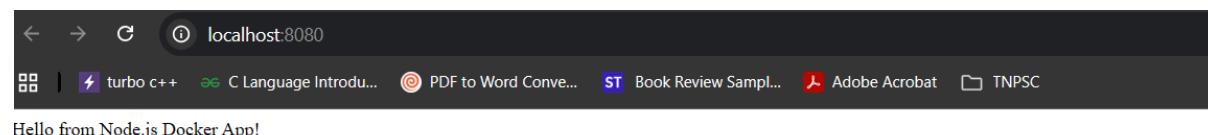
Run in the Powershell.

```
docker build -t regapp-node:1.0 .
docker run -p 8080:8080 regapp-node:1.0
```

- EXPOSE → tells Docker which port the container listens on.
- CMD → default command executed when the container starts.
- docker run -p → maps container port to host port.

```
C:\Users\M.DevaShree\Desktop\regapp-node>docker build -t regapp-node:1.0 .
[+] Building 96.9s (11/11) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 386B
=> [internal] load metadata for docker.io/library/node:20
=> [auth] library/node:pull token for registry-1.docker.io
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [1/5] FROM docker.io/library/node:20@sha256:47dacd49500971c0fbe602323b2d04f6df40a933b123889636fc1f76bf69f58a
=> => resolve docker.io/library/node:20@sha256:47dacd49500971c0fbe602323b2d04f6df40a933b123889636fc1f76bf69f58a
=> => sha256:de002888bed8164550e7a9de5858c8940ba5f45c133a7ba12b83bb4efba51dd1 1.25MB / 1.25MB
=> => sha256:f6819020f2779f375960349b3422c5fe0fe18dc9e80bd6a6f7b74f032b88513 48.41MB / 48.41MB
=> => sha256:fb9baa9d1d1df57d8d063960d7beea8198b17a719381ed80532d617d3d302fa7 3.32KB / 3.32KB
=> => sha256:32885a2b0a589e832bf6b250bd35a528b268360f166af2cd7094d3a14993fc1 211.45MB / 211.45MB
=> => sha256:b82a1e14a32dee2b4a701dc4bee2e6a33e5728a76465e71d9be67bc9d3954cc4 448B / 448B
=> => sha256:2123190679e81d983648da92f1bb9ddc74383512edb00ad64f93d24d00d8807a 64.40MB / 64.40MB
```

```
>docker run -p 8080:8080 regapp-node:1.0
Server running on port 8080
```



Step 4: Push Docker Image to Docker Hub

- Docker Hub is a **public container registry**.
- Pushing images allows any environment (like Kubernetes) to pull and run them.

docker login

```
docker tag regapp-node:1.0 <your-dockerhub-username>/regapp-node:1.0
docker push <your-dockerhub-username>/regapp-node:1.0
```

```
>docker login
Authenticating with existing credentials... [Username: shree2707]

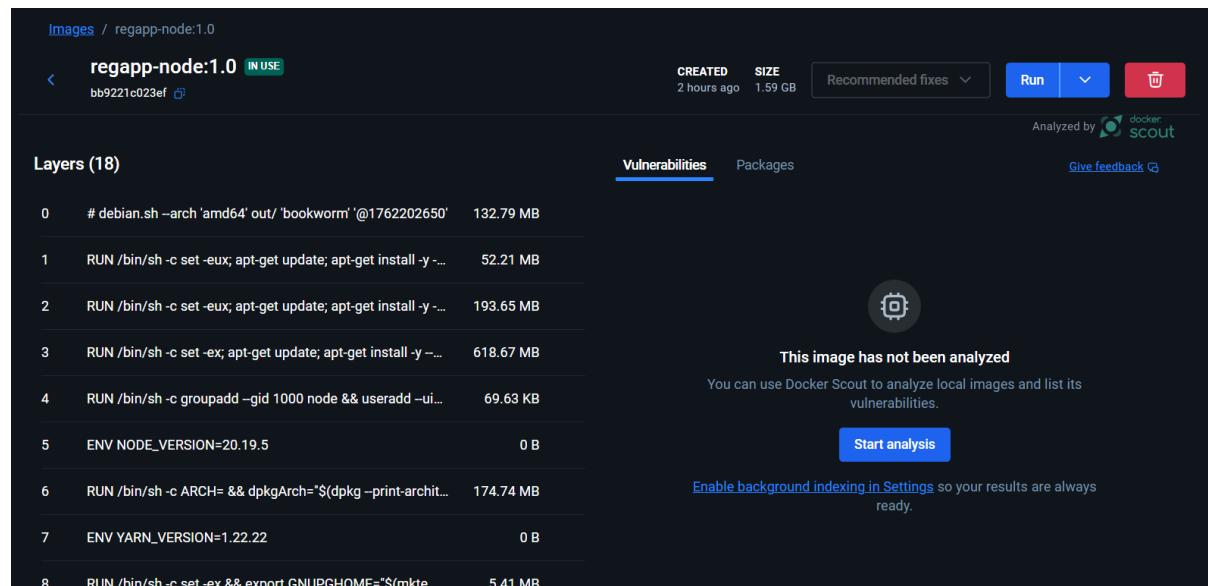
Info → To login with a different account, run 'docker logout' followed by 'docker login'

Login Succeeded

>docker tag regapp-node:1.0 shree2707/regapp-node:1.0

>docker push shree2707/regapp-node:1.0
The push refers to repository [docker.io/shree2707/regapp-node]
77f02c7d2a3e: Pushed
b1ddf1a9b0b4: Pushed
398d1dab3af0: Pushed
32885a2b0a58: Pushed
bb445e472b1b: Pushed
de002888bed8: Pushed
b82a1e14a32d: Pushed
db41a835c7ac: Pushed
5d93aea69798: Pushed
2123190679e8: Pushed
fb9baa9d1d1d: Pushed
e6819020f277: Pushed
508516784b29: Pushed
1.0: digest: sha256:bb9221c023efaad4bd73d1e146e7c6d5a9e27903eed67280504fd17981ea2607 size: 856
```

- Verify image on Docker Hub.



Step 5: Kubernetes Deployment

- Deployment: defines pods (replicas of containers) and manages rolling updates.
- Service: exposes pods to other pods, external network, or both.
- LoadBalancer: allows external access (Docker Desktop assigns a local IP or use port-forward).

Deployment YAML (regapp-node-deploy.yml):

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: regapp-node-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: regapp-node
  template:
    metadata:
      labels:
        app: regapp-node
    spec:
      containers:
        - name: regapp-node
          image: <your-dockerhub-username>/regapp-node:1.0
          ports:
            - containerPort: 8080
```

Service YAML (regapp-node-service.yml):

```
apiVersion: v1
kind: Service
metadata:
  name: regapp-node-service
spec:
  selector:
    app: regapp-node
  ports:
    - port: 8080
      targetPort: 8080
  type: LoadBalancer
```

Step 6: Apply Kubernetes Manifests

Run in the Powershell.

```
kubectl config use-context docker-desktop  
kubectl apply -f regapp-node-deploy.yml  
kubectl apply -f regapp-node-service.yml  
kubectl get pods  
kubectl get svc
```

- `kubectl apply` → applies the configuration in YAML.
- `kubectl get pods` → shows running pods.
- `kubectl get svc` → shows services and external access info.

```
>kubectl apply -f regapp-node-deploy.yml  
deployment.apps/regapp-node-deployment configured  
  
>kubectl apply -f regapp-node-service.yml  
service/regapp-node-service unchanged  
  
>kubectl get pods  
NAME READY STATUS RESTARTS AGE  
regapp-node-deployment-95488b6db-js6fs 1/1 Running 0 15s  
regapp-node-deployment-95488b6db-mp7xv 1/1 Running 0 19s  
regapp-node-deployment-95488b6db-v64tq 1/1 Running 0 24s  
  
>kubectl get svc  
NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE  
kubernetes ClusterIP 10.96.0.1 <none> 443/TCP 5m41s  
regapp-node-service LoadBalancer 10.108.62.74 <pending> 8080:31923/TCP 3m12s
```

Step 7: Access Node.js App

Port-forward:

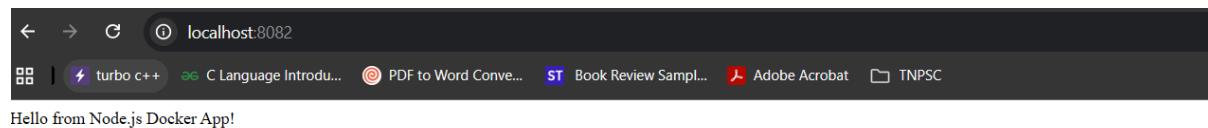
`kubectl port-forward service/regapp-node-service 8082:8080`

Open <http://localhost:8082>

verify the app is running.

- Port-forwarding is useful when LoadBalancer is not available locally.
- Pods are ephemeral; if a pod dies, Kubernetes automatically restarts it.

```
>kubectl port-forward service/regapp-node-service 8082:8080  
Forwarding from 127.0.0.1:8082 -> 8080  
Forwarding from [::1]:8082 -> 8080
```



Step 8: Scale Deployment

Run in the Powershell.

```
kubectl scale deployment regapp-node-deployment --replicas=5  
kubectl get pods
```

- Kubernetes scales pods automatically.
- ReplicaSets ensure the desired number of pods are running at all times.

```
>kubectl scale deployment regapp-node-deployment --replicas=5  
deployment.apps/regapp-node-deployment scaled  
  
>kubectl get pods  
NAME                      READY   STATUS    RESTARTS   AGE  
regapp-node-deployment-95488b6db-8nlml  1/1     Running   0          12s  
regapp-node-deployment-95488b6db-js6fs  1/1     Running   0          2m54s  
regapp-node-deployment-95488b6db-mp7xv  1/1     Running   0          2m58s  
regapp-node-deployment-95488b6db-v64tq  1/1     Running   0          3m3s  
regapp-node-deployment-95488b6db-z2mjx  1/1     Running   0          12s
```

Step 9: Integrate CI/CD

The CI/CD pipeline operates as follows:

1. The developer writes or updates code and pushes it to the **main** branch of GitHub.
2. GitHub Actions automatically detects the change and triggers a workflow.
3. The workflow performs:
 - Code checkout
 - Environment setup
 - Dependency installation
 - Automatic build or test execution
4. If all steps succeed, the build is considered **successful**.

This ensures that the application is always in a **deployable and working state**.

GitHub Actions CI Workflow

The CI workflow is defined in: .github/workflows/ci.yml

Workflow File Contents

```
name: CI Pipeline  
  
on:  
  push:  
    branches: ["main"]  
  pull_request:  
    branches: ["main"]  
  
jobs:
```

```

build:
  runs-on: ubuntu-latest

steps:
  - name: Checkout Repository
    uses: actions/checkout@v3

  - name: Set up Node
    uses: actions/setup-node@v3
    with:
      node-version: 18

  - name: Install Dependencies
    run: npm install

  - name: Run Tests (Optional)
    run: npm test || echo "No tests available"

  - name: Build Application (Optional)
    run: npm run build || echo "No build step configured"

```

GitHub Actions CD Workflow

The CD workflow is defined in: .github/workflows/cd.yml

Workflow File Contents

```

name: CD to Docker Desktop Kubernetes

on:
  push:
    branches: ["main"]

jobs:
  deploy:
    runs-on: ubuntu-latest

steps:
  - name: Checkout code
    uses: actions/checkout@v3

  - name: Login to Docker Hub
    run: echo "${{ secrets.DOCKER_PASSWORD }}" | docker login -u "${{ secrets.DOCKER_USERNAME }}" --password-stdin

  - name: Build Docker Image
    run: docker build -t ${{ secrets.DOCKER_USERNAME }}/regapp-node:latest .

  - name: Push Docker Image
    run: docker push ${{ secrets.DOCKER_USERNAME }}/regapp-node:latest

```

```

- name: Set up kubectl
  uses: azure/setup-kubectl@v3

- name: Configure kubeconfig
  run: |
    mkdir -p ~/.kube
    echo "${{ secrets.KUBE_CONFIG }}" > ~/.kube/config

- name: Deploy to Kubernetes
  run: |
    kubectl set image deployment/regapp-node-deployment regapp-node={{ secrets.DOCKER_USERNAME }}/regapp-node:latest
    kubectl rollout status deployment/regapp-node-deployment

```

Summary

Triggered via push now
Shree2722 pushed → `main`

Status: Success Total duration: 17s Artifacts: -

`ci.yml` on push

`build` 17s

Summary

Triggered via push 52 minutes ago
Shree2722 pushed → `main`

Status: Success Total duration: 15s Artifacts: -

`ci.yml` on push

`build` 15s

Step 10: Verification

Run in the Powershell.

```

kubectl get pods -o wide
kubectl get svc -o wide
kubectl describe deployment regapp-node-deployment

```

```

>kubectl get pods -o wide
NAME           READY   STATUS    RESTARTS   AGE     IP          NODE      NOMINATED NODE   READINESS GATES
regapp-node-deployment-95488b6db-8nlml   1/1    Running   0          175m   10.1.0.12   docker-desktop   <none>        <none>
regapp-node-deployment-95488b6db-j56fs   1/1    Running   0          178m   10.1.0.11   docker-desktop   <none>        <none>
regapp-node-deployment-95488b6db-mp7xv   1/1    Running   0          178m   10.1.0.10   docker-desktop   <none>        <none>
regapp-node-deployment-95488b6db-v64tq   1/1    Running   0          178m   10.1.0.9    docker-desktop   <none>        <none>
regapp-node-deployment-95488b6db-z2mjx   1/1    Running   0          175m   10.1.0.13   docker-desktop   <none>        <none>

>kubectl get svc -o wide
NAME           TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)        AGE     SELECTOR
kubernetes     ClusterIP   10.96.0.1   <none>        443/TCP       3h3m   <none>
regapp-node-service   LoadBalancer  10.108.62.74  <pending>   8080:31923/TCP 3h1m   app=regapp-node

```

```
>kubectl describe deployment regapp-node-deployment
Name:           regapp-node-deployment
Namespace:      default
CreationTimestamp:  Wed, 05 Nov 2025 17:49:15 +0530
Labels:         <none>
Annotations:   deployment.kubernetes.io/revision: 2
Selector:       app=regapp-node
Replicas:       5 desired | 5 updated | 5 total | 5 available | 0 unavailable
StrategyType:  RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=regapp-node
  Containers:
    regapp-node:
      Image:      shree2707/regapp-node:1.0
      Port:       8080/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
      Node-Selectors: <none>
      Tolerations:  <none>
  Conditions:
    Type        Status  Reason
    ----        ----   -----
    Progressing  True    NewReplicaSetAvailable
    Available   True    MinimumReplicasAvailable
OldReplicaSets: regapp-node-deployment-6f695488bf (0/0 replicas created)
NewReplicaSet:  regapp-node-deployment-95488b6db (5/5 replicas created)
Events:        <none>
```

Preparation	20	
Implementation	20	
Viva	15	
Output	10	
Record	10	
Total	75	

Result:

Thus ,the application was successfully containerized using a Dockerfile, which packaged the source code along with all necessary dependencies into a portable container image.

Aim:

To Develop an ETL process using SQL scripts to extract data from a source system (e.g., transactional database), transform it (e.g., aggregate, cleanse, join), and load it into a data warehouse.

1. Open MySQL

On Windows CMD or Terminal

```
mysql -u root -p
```

Then enter your MySQL password. After login, you'll see:

```
mysql>
```

2. Create Source Database (Transactional System)

Create the database that represents the retail transactions.

```
CREATE DATABASE retail_source;
```

```
USE retail_source;
```

```
mysql> CREATE DATABASE retail_source;
Query OK, 1 row affected (0.01 sec)
```

```
mysql> USE retail_source;
Database changed
```

3. Create Source Table

This table stores raw sales transactions.

```
CREATE TABLE sales_transaction (
    Transaction_ID INT PRIMARY KEY,
    Date DATE,
    Customer_ID INT,
    Gender VARCHAR(10),
    Age INT,
    Product_Category VARCHAR(50),
    Quantity INT,
    Price_per_Unit DECIMAL(10,2),
    Total_Amount DECIMAL(10,2)
);
```

```
mysql> CREATE TABLE sales_transaction (Transaction_ID INT PRIMARY KEY, Date DATE, Customer_ID INT, Gender VARCHAR(10), Age INT, Product_Category VARCHAR(50), Q
uantity INT, Price_per_Unit DECIMAL(10,2), Total_Amount DECIMAL(10,2));
Query OK, 0 rows affected (0.04 sec)
```

4. Insert Data

Insert a larger dataset to make ETL.

Sample Data:

```
INSERT INTO sales_transaction VALUES
```

```
(1, '2025-01-01', 101, 'Male', 275, 'Electronics', 2, 500.00, 1000.00),
(2, '2025-01-02', 102, 'Female', 30, 'Grocery', 5, 50.00, 250.00),
(3, '2025-01-02', 103, 'Male', 40, 'Clothing', 3, 200.00, 600.00),
(4, '2025-01-03', 104, 'Female', 22, 'Furniture', 1, 800.00, 800.00),
(5, '2025-01-03', 105, 'Male', 35, 'Clothing', 4, 150.00, 600.00),
(6, '2025-01-04', 106, 'Female', 28, 'Electronics', 1, 700.00, 700.00),
(7, '2025-01-04', 107, 'Male', 42, 'Grocery', 10, 40.00, 400.00),
(8, '2025-01-05', 108, 'Female', 33, 'Furniture', 2, 1200.00, 2400.00),
(9, '2025-01-05', 109, 'Male', 27, 'Beauty', 5, 100.00, 500.00),
(10, '2025-01-06', 110, 'Female', 38, 'Electronics', 1, 900.00, 900.00),
(11, '2025-01-07', 111, 'Male', 29, 'Furniture', 1, 1500.00, 1500.00),
(12, '2025-01-07', 112, 'Female', 24, 'Beauty', 3, 100.00, 300.00),
(13, '2025-01-08', 113, 'Male', 31, 'Beauty', 2, 200.00, 400.00),
(14, '2025-01-08', 114, 'Female', 41, 'Electronics', 2, 850.00, 1700.00),
(15, '2025-01-09', 115, 'Male', 37, 'Grocery', 8, 60.00, 480.00),
(16, '2025-01-09', 116, 'Female', 32, 'Clothing', 3, 250.00, 750.00),
(17, '2025-01-10', 117, 'Male', 45, 'Electronics', 1, 1000.00, 1000.00),
(18, '2025-01-10', 118, 'Female', 29, 'Furniture', 1, 900.00, 900.00),
(19, '2025-01-11', 119, 'Male', 34, 'Grocery', 5, 55.00, 275.00),
(20, '2025-01-12', 120, 'Female', 26, 'Clothing', 2, 200.00, 400.00);
```

```
mysql> INSERT INTO sales_transaction VALUES(1, '2025-01-01', 101, 'Male', 275, 'Electronics', 2, '2025-01-02', 102, 'Female', 30, 'Grocery', 5, 50.00, 250.00),(3, '2025-01-02', 103, 'Male', 40, 'Clothing', 3, 200.00, 600.00),(4, '2025-01-03', 104, 'Female', 22, 'Furniture', 1, 800.00, 800.00),(5, '2025-01-03', 105, 'Male', 35, 'Clothing', 4, 150.00, 600.00),(6, '2025-01-04', 106, 'Female', 28, 'Electronics', 1, 700.00, 700.00),(7, '2025-01-04', 107, 'Male', 42, 'Grocery', 10, 40.00, 400.00),(8, '2025-01-05', 108, 'Female', 33, 'Furniture', 2, 1200.00, 2400.00),(9, '2025-01-05', 109, 'Male', 27, 'Beauty', 5, 100.00, 500.00),(10, '2025-01-06', 110, 'Female', 38, 'Electronics', 1, 900.00, 900.00),(11, '2025-01-07', 111, 'Male', 29, 'Furniture', 1, 1500.00, 1500.00),(12, '2025-01-07', 112, 'Female', 24, 'Beauty', 3, 100.00, 300.00),(13, '2025-01-08', 113, 'Male', 31, 'Beauty', 2, 200.00, 400.00),(14, '2025-01-08', 114, 'Female', 41, 'Electronics', 2, 850.00, 1700.00),(15, '2025-01-09', 115, 'Male', 37, 'Grocery', 8, 60.00, 480.00),(16, '2025-01-09', 116, 'Female', 32, 'Clothing', 3, 250.00, 750.00),(17, '2025-01-10', 117, 'Male', 45, 'Electronics', 1, 1000.00, 1000.00),(18, '2025-01-10', 118, 'Female', 29, 'Furniture', 1, 900.00, 900.00),(19, '2025-01-11', 119, 'Male', 34, 'Grocery', 5, 55.00, 275.00),(20, '2025-01-12', 120, 'Female', 26, 'Clothing', 2, 200.00, 400.00);
Query OK, 20 rows affected (0.02 sec)
Records: 20  Duplicates: 0  Warnings: 0
```

5. Create Data Warehouse Database

This is where **transformed data** will be stored.

```
CREATE DATABASE retail_dw;
```

```
USE retail_dw;
```

```
mysql> CREATE DATABASE retail_dw;
Query OK, 1 row affected (0.01 sec)
```

```
mysql> USE retail_dw;
Database changed
```

6. Create Dimension Tables

1. Customer Dimension

```
CREATE TABLE dim_customer (
    Customer_ID INT PRIMARY KEY,
    Gender VARCHAR(10),
    Age INT
);
```

2. Product Dimension

```
CREATE TABLE dim_product (
    Product_ID INT AUTO_INCREMENT PRIMARY KEY,
    Product_Category VARCHAR(50)
);
```

3. Time Dimension

```
CREATE TABLE dim_time (
    Date DATE PRIMARY KEY,
    Year INT,
    Month INT,
    Day INT
);
```

```
mysql> CREATE TABLE dim_customer (Customer_ID INT PRIMARY KEY,Gender VARCHAR(10),Age INT);
Query OK, 0 rows affected (0.04 sec)

mysql> CREATE TABLE dim_product (Product_ID INT AUTO_INCREMENT PRIMARY KEY,Product_Category VARCHAR(50));
Query OK, 0 rows affected (0.03 sec)

mysql> CREATE TABLE dim_time (Date DATE PRIMARY KEY,Year INT,Month INT,Day INT);
Query OK, 0 rows affected (0.04 sec)
```

7. Create Fact Table

This stores the measurable data the sales transactions.

```
CREATE TABLE fact_sales (
    Transaction_ID INT PRIMARY KEY,
    Date DATE,
    Customer_ID INT,
    Product_ID INT,
    Quantity INT,
    Total_Amount DECIMAL(10,2),
    FOREIGN KEY (Customer_ID) REFERENCES dim_customer(Customer_ID),
    FOREIGN KEY (Product_ID) REFERENCES dim_product(Product_ID),
    FOREIGN KEY (Date) REFERENCES dim_time(Date));
```

```
mysql> CREATE TABLE fact_sales (Transaction_ID INT PRIMARY KEY, Date DATE, Customer_ID INT, Product_ID INT, Quantity INT, Total_Amount DECIMAL(10,2), FOREIGN KEY (Customer_ID) REFERENCES dim_customer(Customer_ID), FOREIGN KEY (Product_ID) REFERENCES dim_product(Product_ID), FOREIGN KEY (Date) REFERENCES dim_time(Date))
;
Query OK, 0 rows affected (0.09 sec)
```

8. Extract Data

Move raw data into staging tables before cleaning.

```
CREATE TABLE staging_sales AS
SELECT * FROM retail_source.sales_transaction;
```

```
mysql> CREATE TABLE staging_sales AS SELECT * FROM retail_source.sales_transaction;
Query OK, 20 rows affected (0.03 sec)
Records: 20  Duplicates: 0  Warnings: 0
```

9. Transform Data

-> Remove duplicates

```
DELETE s1 FROM staging_sales s1
JOIN staging_sales s2
WHERE s1.Transaction_ID > s2.Transaction_ID
AND s1.Transaction_ID = s2.Transaction_ID;
```

```
mysql> DELETE s1 FROM staging_sales s1 JOIN staging_sales s2 WHERE s1.Transaction_ID > s2.Transaction_ID AND s1.Transaction_ID = s2.Transaction_ID;
Query OK, 0 rows affected (0.01 sec)
```

-> Handle missing or invalid values

```
DELETE FROM staging_sales
WHERE Total_Amount IS NULL OR Quantity IS NULL;
```

```
mysql> DELETE FROM staging_sales WHERE Total_Amount IS NULL OR Quantity IS NULL;
Query OK, 0 rows affected (0.00 sec)
```

-> Check inconsistencies

```
UPDATE staging_sales
SET Total_Amount = Quantity * Price_per_Unit
WHERE Total_Amount != Quantity * Price_per_Unit;
```

```
mysql> UPDATE staging_sales SET Total_Amount = Quantity * Price_per_Unit WHERE Total_Amount != Quantity * Price_per_Unit;
Query OK, 0 rows affected (0.00 sec)
Rows matched: 0  Changed: 0  Warnings: 0
```

10. Load Data into Dimensions

Customer

```
INSERT INTO dim_customer (Customer_ID, Gender, Age)
SELECT DISTINCT Customer_ID, Gender, Age FROM staging_sales;
```

```
mysql> INSERT INTO dim_customer (Customer_ID, Gender, Age) SELECT DISTINCT Customer_ID, Gender, Age FROM staging_sales;
Query OK, 20 rows affected (0.01 sec)
Records: 20  Duplicates: 0  Warnings: 0
```

Product

INSERT INTO dim_product (Product_Category)

SELECT DISTINCT Product_Category FROM staging_sales;

```
mysql> INSERT INTO dim_product (Product_Category) SELECT DISTINCT Product_Category FROM staging_sales;
Query OK, 5 rows affected (0.01 sec)
Records: 5  Duplicates: 0  Warnings: 0
```

Time

INSERT INTO dim_time (Date, Year, Month, Day)

SELECT DISTINCT

Date,

YEAR(Date),

MONTH(Date),

DAY(Date)

FROM staging_sales;

```
mysql> INSERT INTO dim_time (Date, Year, Month, Day) SELECT DISTINCT Date, YEAR(Date), MONTH(Date), DAY(Date) FROM staging_sales;
Query OK, 12 rows affected (0.01 sec)
Records: 12  Duplicates: 0  Warnings: 0
```

11. LOAD Data into Fact Table

Join the product dimension to get Product_ID.

INSERT INTO fact_sales (Transaction_ID, Date, Customer_ID, Product_ID, Quantity, Total_Amount)

SELECT

s.Transaction_ID,

s.Date,

s.Customer_ID,

p.Product_ID,

s.Quantity,

s.Total_Amount

FROM staging_sales s

JOIN dim_product p

ON s.Product_Category = p.Product_Category;

```
mysql> INSERT INTO fact_sales (Transaction_ID, Date, Customer_ID, Product_ID, Quantity, Total_Amount) SELECT s.Transaction_ID, s.Date, s.Customer_ID, p.Product_ID, s.Quantity, s.Total_Amount FROM staging_sales s JOIN dim_product p ON s.Product_Category = p.Product_Category;
Query OK, 20 rows affected (0.01 sec)
Records: 20  Duplicates: 0  Warnings: 0
```

12. Data Quality Check

Count checks

SELECT COUNT() FROM fact_sales;*

```
mysql> SELECT COUNT(*) FROM fact_sales;
+-----+
| COUNT(*) |
+-----+
|      20 |
+-----+
1 row in set (0.01 sec)
```

SELECT COUNT(DISTINCT Customer_ID) FROM dim_customer;

```
mysql> SELECT COUNT(DISTINCT Customer_ID) FROM dim_customer;
+-----+
| COUNT(DISTINCT Customer_ID) |
+-----+
|          20 |
+-----+
1 row in set (0.00 sec)
```

Null checks

*SELECT * FROM fact_sales WHERE Total_Amount IS NULL;*

```
mysql> SELECT * FROM fact_sales WHERE Total_Amount IS NULL;
Empty set (0.00 sec)
```

13. Reporting Queries

1. Total Sales Revenue by Product Category

*SELECT p.Product_Category, SUM(f.Total_Amount) AS Total_Sales
FROM fact_sales f
JOIN dim_product p ON f.Product_ID = p.Product_ID
GROUP BY p.Product_Category;*

```
mysql> SELECT p.Product_Category, SUM(f.Total_Amount) AS Total_Sales FROM fact_sales f JOIN dim_product p ON f.Product_ID = p.Product_ID GROUP BY p.Product_Category;
+-----+-----+
| Product_Category | Total_Sales |
+-----+-----+
| Electronics      |    1000.00 |
| Grocery          |     250.00 |
| Clothing          |    1200.00 |
| Furniture         |     800.00 |
+-----+-----+
4 rows in set (0.00 sec)
```

2. Customer Demographics (Average Spending by Gender)

*SELECT c.Gender, AVG(f.Total_Amount) AS Avg_Spending
FROM fact_sales f
JOIN dim_customer c ON f.Customer_ID = c.Customer_ID
GROUP BY c.Gender;*

```

mysql> SELECT c.Gender, AVG(f.Total_Amount) AS Avg_Spending
-> FROM fact_sales f
-> JOIN dim_customer c ON f.Customer_ID = c.Customer_ID
-> GROUP BY c.Gender;
+-----+
| Gender | Avg_Spending |
+-----+
| Male   |    733.33333 |
| Female |    525.00000 |
+-----+
2 rows in set (0.00 sec)

```

3. Sales Trend Over Time

```

SELECT t.Month, SUM(f.Total_Amount) AS Monthly_Sales
FROM fact_sales f
JOIN dim_time t ON f.Date = t.Date
GROUP BY t.Month
ORDER BY t.Month;

```

```

mysql> SELECT t.Month, SUM(f.Total_Amount) AS Monthly_Sales
-> FROM fact_sales f
-> JOIN dim_time t ON f.Date = t.Date
-> GROUP BY t.Month
-> ORDER BY t.Month;
+-----+
| Month | Monthly_Sales |
+-----+
| 1      |      3250.00 |
+-----+
1 row in set (0.00 sec)

```

Preparation	20	
Implementation	20	
Viva	15	
Output	10	
Record	10	
Total	75	

Result:

Thus, the ETL process successfully extracted raw transactional data, cleansed and transformed it, and loaded it into the data warehouse.